Maksym Andriushchenko

CURRICULUM VITAE

andriushchenko.me · maksym@andriushchenko.me · Google Scholar · GitHub

Education

École Polytechnique Fédérale de Lausanne (EPFL), Switzerland *(Sep 2019 – Oct 2024)*
PhD Degree in Computer Science advised by Nicolas Flammarion
Thesis title: "Understanding Generalization and Robustness in Modern Deep Learning"
Committee: Zico Kolter, Sebastien Bubeck, Rachid Guerraoui, Florent Krzakala

Saarland University, Germany *(Oct 2016 – Aug 2019)*
Master's Degree in Computer Science advised by Matthias Hein from the University of Tübingen
Thesis title: "Provable Adversarial Defenses for Boosting"

Dnipropetrovsk National University of Railway Transport, Ukraine *(Sep 2012 – Jun 2016)*
Bachelor's Degree in Software Engineering — *with honors*

Work

ELLIS Institute and Max Planck Institute	**Time**: September 2025 – now **Role**: Principal investigator at the ELLIS Institute Tübingen and independent group leader at the Max Planck Institute for Intelligent Systems
EPFL	**Time**: September 2019 – August 2025 **Role**: PhD student → postdoctoral researcher supervised by Nicolas Flammarion
Gray Swan AI	**Time**: February 2024 – November 2024 (part-time consultant) **Role**: Member of technical staff working with leading AI safety organizations (Anthropic, UK AI Safety Institute, Center for AI Safety)
Adobe Research	**Time**: July 2021 – October 2021 **Role**: Research internship supervised by John Collomosse
PrivatBank	**Time**: November 2015 – June 2016 (part-time) **Role**: Data scientist working on predictive modeling and e-commerce applications

Awards and Grants

Independent grants as a PI or co-PI	'The Challenge of Test Awareness in AI Safety Evaluation', The Alignment Project from the UK AISI ($882,000, co-PI) 'Promoting Cooperation and Strategic Alignment in Multi-Agent LLMs', Large Grants from Swiss AI ($68,000 and 220,000 GPU hours, co-PI) 'Faithfulness and Safety of Latent Chain-of-Thought Reasoning', Safety in the Inference-Time Compute Paradigm from Schmidt Sciences ($497,000, co-PI) Academic Start-up Package, Technical AI Safety Program from Open Philanthropy ($1,000,000, single PI)
Co-written grants	AI Safety Science grant ($200k funded by Schmidt Sciences in 2025) Safe GenAI via Robust Content Moderation Models ($100k funded by Google in 2024) Safety, Robustness, and Alignment of LLM agents ($40k funded by Google in 2024) Google Research Collab 2022-2023 ($80k for one year + $20k in cloud compute)

Main awards	ELLIS PhD Award ($2.5k award) Patrick Denantes Memorial Prize for the best PhD thesis in the School of Computer and Communication Sciences at EPFL ($5k award) G-Research PhD Thesis Prize in Quantitative Research (2nd place, $3k award) Google PhD fellowship 2022-2025 ($80k per year for 3 years) Open Philanthropy AI PhD Fellowship 2022-2024 ($10k per year for travel/equipment) EDIC PhD fellowship from EPFL for the first year ($60k) DAAD MSc scholarship for 2 years to study at Saarland University ($20k)
Awards for papers and competitions	Trojan detection challenge at SatML'24: **first place** Swiss AI Safety Prize (2024): **award for one of the top paper submissions** Joint Conference of Korean AI Association (2023): **best paper award** ICLR Workshop on Security & Safety in ML Systems (2021): **best paper honorable mention** Swiss Machine Learning Day (2019): **best paper award**

Press Coverage

2025	Forbes: Embedding LLM Circuit Breakers Into AI Might Save Us From A Whole Lot Of Ghastly Troubles Le Figaro: When ChatGPT Goes Off the Rails, Should We Be Worried? 24 heures: When ChatGPT Goes Off the Rails, Should We Be Worried?
2024	SwissInfo: Swiss researchers find security flaws in AI models Le Temps: Bomb making, cannibalism or cyber attacks: it is always possible to ask AI for advice on explosive topics Blick: EPFL discovers security flaws in several AI models TechXplore: Can we convince AI to answer harmful requests? Léman Bleu: EPFL: Security flaws in AI models ActuIA: EPFL study: the limits of LLMs in the face of adaptive attacks

Academic Service

Participant	International AI Safety Report (Second Edition): Chapter Lead (July 2025 – February 2026) Wilson Center Fellowship Program: A Path Toward AI Policy (November 2024 – January 2025) Red teaming of the OpenAI fine-tuning service as an external expert (October 2023) Robust AI 4-day workshop organized by AirBus AI Research and TNO (January 2021)
Area chair	ICML'26
Reviewer	PNAS, TMLR, ICLR'25, NeurIPS'24, NeurIPS'23, ICML'23, NeurIPS'22 (**top reviewer**), ICML'22, NeurIPS'21, ICML'21, CVPR'21, ICLR'21 (**outstanding reviewer**), NeurIPS'20 (**top 10% reviewers**)

Program committee in workshops	**NeurIPS'24** Red Teaming GenAI Workshop, **NeurIPS'24** SATA Workshop: Towards Safe & Trustworthy Agents, **NeurIPS'24** 3rd Workshop on New Frontiers in Adversarial ML, **ICML'24** Workshop on the Next Generation of AI Safety, **NeurIPS'23** R0-FoMo Workshop on Robustness of Few-shot and Zero-shot Learning in Foundation Models, **NeurIPS'23** Workshop on Distribution Shifts: New Frontiers with Foundation Models, **ICML'23** 2nd ICML Workshop on New Frontiers in Adversarial ML, **ICLR'23** Workshop on Pitfalls of Limited Data and Computation for Trustworthy ML, **NeurIPS'22** Workshop on Distribution Shifts, **NeurIPS'22** ML Safety Workshop, **ICML'22** New Frontiers in Adversarial ML, **ICML'22** Principles of Distribution Shift, **NeurIPS'21** Distribution Shifts: Connecting Methods and Applications, **ICML'21** Uncertainty and Robustness in Deep Learning, **CVPR'21** Adversarial ML in Real-World Computer Vision Systems, **ICLR'21** Robust and Reliable ML in the Real World, **ICLR'21** Security and Safety in ML Systems, **ICML'20** Uncertainty and Robustness in Deep Learning, **CVPR'20** Adversarial ML in Computer Vision, **ICLR'20** Towards Trustworthy ML (**best reviewer award**)
Outreach activities	National coordinator for Switzerland at #ScienceForUkraine Coordinator for Switzerland and admission officer at the Ukrainian Global University AI lectures at a summer camp for displaced Ukrainian children in Romania

Current Group

PhD students	Ben Rank, David Schmotz, Jeremy Qin, Jeanne Salle (co-supervised with Sahar Abdelnabi), Alexander Panfilov (co-supervised with Jonas Geiping), Hardik Bhatnagar (co-supervised with Matthias Bethge)
Research interns	Yuchen Zhang, Jehyeok Yeon, Anietta Weckauff, Raffaele Mura
Master's thesis students	Jonas Wiedermann-Möller, Changling Li, Derck Prinzhorn, Lena Libon

Past Students

Agatha Duzan (EPFL)	**MSc project (2025)**: A Benchmark for Measuring Safety of Computer Use Agents
Thomas Kuntz (EPFL)	**MSc project (2025)**: A Benchmark for Measuring Safety of Computer Use Agents
Joshua Freeman (ETH Zürich)	**MSc project (2024)**: Exploring Memorization and Copyright Violation in Frontier Large Language Models (*accepted at NeurIPS 2024 Safe Generative AI Workshop*)
Hao Zhao (EPFL)	**MSc thesis (2023)**: Long Is More for Alignment: A Simple but Tough-to-Beat Baseline for Instruction Fine-Tuning (*published at ICML'24, nominated for EPFL Outstanding Master's Thesis, featured in* MIT Technology Review China)
Hichem Hadhri (EPFL)	**MSc project (2023)**: Understanding overfitting in large language models
Tiberiu Musat (EPFL)	**BSc project (2023)**: Investigating key components for fast optimization of deep networks
Francesco d'Angelo (EPFL)	**PhD semester project (2023)**: Understanding the role of weight decay in deep learning (*published at NeurIPS'24*)
Théau Vannier (EPFL)	**MSc project (2023)**: Understanding the training instability of transformers
Joshua Freeman (EPFL)	**BSc project (2022, unofficial)**: Automatic recognition of unexploded ordnance using transfer learning
Jana Vuckovic (EPFL)	**MSc project (2022)**: Rethinking the relationship between sharpness and generalization (*follow-up work is published at ICML'23*)
Mehrdad Saberi (EPFL)	**Summer internship (2021)**: Wasserstein adversarial training and perceptual robustness
Edoardo Debenedetti (EPFL)	**MSc project (2021)**: RobustBench: a standardized adversarial robustness benchmark (*published at NeurIPS'21 Datasets and Benchmarks Track; Best Paper Honorable Mention Prize at* ICLR 2021 Workshop on Security and Safety in ML Systems)

Klim Kireev (EPFL)	**PhD semester project (2020)**: On the effectiveness of adversarial training against common corruptions (*published at UAI'22*)
Etienne Bonvin (EPFL)	**MSc project (2020)**: Adversarial robustness of kernel methods
Oriol Barbany (EPFL)	**MSc project (2019)**: Affine-invariant robust training (co-advised with Sebastian Stich)

Teaching Experience

EPFL	**Probability & Statistics 2021, 2022** (by E. Abbé), **Machine Learning 2020, 2021, 2022, 2023** (by M. Jaggi, N. Flammarion), **Advanced Algorithms 2020** (by M. Kapralov)
MPI for Informatics	**Machine Learning 2018-2019** (lecturer: B. Schiele)
Saarland University	**Neural Networks: Implementation and Application 2017** (lecturer: D. Klakow)

Selected Publications

B. Rank, H. Bhatnagar, A. Prabhu, S. Eisenberg, K. Nguyen, M. Bethge, **M. Andriushchenko**. *PostTrainBench: Can LLM Agents Automate LLM Post-Training?* (arXiv, March 2026)

M. Andriushchenko, A. Souly, M. Dziemian, D. Duenas, M. Lin, J. Wang, D. Hendrycks, A. Zou, Z. Kolter, M. Fredrikson, E. Winsor, J. Wynne, Y. Gal, X. Davies. *AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents* (ICLR 2025)

M. Andriushchenko, F. Croce, N. Flammarion. *Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks* (ICLR 2025)

F. Croce*, **M. Andriushchenko***, V. Sehwag*, E. Debenedetti*, N. Flammarion, M. Chiang, P. Mittal, M. Hein. *RobustBench: a standardized adversarial robustness benchmark* (NeurIPS 2021 Datasets and Benchmarks Track, **Best Paper Honorable Mention Prize** at ICLR 2021 Workshop on Security and Safety in Machine Learning Systems)

M. Andriushchenko*, F. Croce*, N. Flammarion, M. Hein. *Square Attack: a Query-Efficient Black-Box Adversarial Attack via Random Search* (ECCV 2020)

Full Publication List

B. Rank, H. Bhatnagar, A. Prabhu, S. Eisenberg, K. Nguyen, M. Bethge, **M. Andriushchenko**. *PostTrainBench: Can LLM Agents Automate LLM Post-Training?* (arXiv, March 2026)

D. Schmotz, L. Beurer-Kellner, S. Abdelnabi, **M. Andriushchenko**. *Skill-Inject: Measuring Agent Vulnerability to Skill File Attacks* (arXiv, February 2026)

N. Talokar, A.K. Tarun, M. Mandal, **M. Andriushchenko**, A. Bosselut. *Helpful to a Fault: Measuring Illicit Assistance in Multi-Turn, Multilingual LLM Agents* (arXiv, February 2026)

D. Fan, S. Delsad, N. Flammarion, **M. Andriushchenko**. *HalluHard: A Hard Multi-Turn Hallucination Benchmark* (arXiv, February 2026)

M. Terekhov, A. Panfilov, D. Dzenhaliou, C. Gulcehre, **M. Andriushchenko**, A. Prabhu, J. Geiping. *Adaptive Attacks on Trusted Monitors Subvert AI Control Protocols* (ICLR 2026)

A. Panfilov, E. Kortukov, K. Nikolić, M. Bethge, S. Lapuschkin, W. Samek, A. Prabhu, **M. Andriushchenko**, J. Geiping. *Strategic Dishonesty Can Undermine AI Safety Evaluations of Frontier LLMs* (ICLR 2026)

D. Schmotz, S. Abdelnabi, **M. Andriushchenko**. *Agent Skills Enable a New Class of Realistic and Trivially Simple Prompt Injections* (arXiv, October 2025)

T. Kuntz, A. Duzan, H. Zhao, F. Croce, Z. Kolter, N. Flammarion, **M. Andriushchenko**. *OS-Harm: A Benchmark for Measuring Safety of Computer Use Agents* (NeurIPS 2025 Datasets and Benchmarks Track, **spotlight**)

C. Yueh-Han, N. Joshi, Y. Chen, **M. Andriushchenko**, R. Angell, H. He. *Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors* (ICLR 2026)

A. Panfilov, P. Kassianik, **M. Andriushchenko**, J. Geiping. *Capability-Based Scaling Laws for LLM Red-Teaming*

(ICLR 2026)

M. Andriushchenko, A. Souly, M. Dziemian, D. Duenas, M. Lin, J. Wang, D. Hendrycks, A. Zou, Z. Kolter, M. Fredrikson, E. Winsor, J. Wynne, Y. Gal, X. Davies. *AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents* (ICLR 2025)

M. Andriushchenko, F. Croce, N. Flammarion. *Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks* (ICLR 2025)

M. Andriushchenko, N. Flammarion. *Does Refusal Training in LLMs Generalize to the Past Tense?* (ICLR 2025)

H. Zhao, **M. Andriushchenko**, F. Croce, N. Flammarion. *Is In-Context Learning Sufficient for Instruction Following in LLMs?* (ICLR 2025)

J. Freeman, C. Rippe, E. Debenedetti, **M. Andriushchenko**. *Exploring Memorization and Copyright Violation in Frontier LLMs: A Study of the New York Times v. OpenAI 2023 lawsuit* (NeurIPS 2024 Safe Generative AI Workshop)

A. Zou, L. Phan, J. Wang, D. Duenas, M. Lin, **M. Andriushchenko**, R. Wang, Z. Kolter, M. Fredrikson, D. Hendrycks. *Improving Alignment and Robustness with Short Circuiting* (NeurIPS 2024)

M. Andriushchenko*, F. D'Angelo*, A. Varre, N. Flammarion. *Why Do We Need Weight Decay in Modern Deep Learning?* (NeurIPS 2024)

P. Chao*, E. Debenedetti*, A. Robey*, **M. Andriushchenko***, F. Croce, V. Sehwag, E. Dobriban, N. Flammarion, G.J. Pappas, F. Tramer, H. Hassani, E. Wong. *JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models* (NeurIPS 2024 Datasets and Benchmarks Track)

J. Rando, F. Croce, K. Mitka, S. Shabalin, **M. Andriushchenko**, N. Flammarion, F. Tramèr. Competition Report: Finding Universal Jailbreak Backdoors in Aligned LLMs (arXiv, April 2024)

H. Zhao, **M. Andriushchenko**, F. Croce, N. Flammarion. *Long Is More for Alignment: A Simple but Tough-to-Beat Baseline for Instruction Fine-Tuning* (ICML 2024)

L. Adilova, **M. Andriushchenko**, M. Kamp, A. Fischer, M. Jaggi. *Layer-Wise Linear Mode Connectivity* (ICLR 2024)

M. Andriushchenko. *Adversarial Attacks on GPT-4 via Simple Random Search* (arXiv, December 2023)

E. Debenedetti, Z. Wan, **M. Andriushchenko**, V. Sehwag, K. Bhardwaj, B. Kailkhura. *Scaling Compute Is Not All You Need for Adversarial Robustness* (ICLR 2024 Workshop on Reliable and Responsible Foundation Models)

S. Shin, D. Lee, **M. Andriushchenko**, N. Lee. *The Effects of Overparameterization on Sharpness-Aware Minimization: An Empirical and Theoretical Analysis* (UAI 2025, **best paper award** at the Joint Conference of Korean Artificial Intelligence Association (2023))

M. Andriushchenko, D. Bahri, H. Mobahi, N. Flammarion. *Sharpness-Aware Minimization Leads to Low-Rank Features* (NeurIPS 2023)

K. Kireev, **M. Andriushchenko**, C. Troncoso, N. Flammarion. *Transferable Adversarial Robustness for Categorical Data via Universal Robust Embeddings* (NeurIPS 2023)

M. Andriushchenko, F. Croce, M. Müller, M. Hein, N. Flammarion. *A modern look at the relationship between sharpness and generalization* (ICML 2023)

M. Andriushchenko, A. Varre, L. Pillaud-Vivien, N. Flammarion. *SGD with large step sizes learns sparse features* (ICML 2023)

K. Kireev*, **M. Andriushchenko***, N. Flammarion. *On the effectiveness of adversarial training against common corruptions* (UAI 2022)

M. Rose, S. Reinsone, **M. Andriushchenko**, M. Bartosiak, A. Bobak et al. #ScienceForUkraine: an Initiative to Support the Ukrainian Academic Community. *"3 Months Since Russia's Invasion of Ukraine", February 26 – May 31, 2022* (SSRN, 2022)

M. Andriushchenko, N. Flammarion. *Towards Understanding Sharpness-Aware Minimization* (ICML 2022)

M. Andriushchenko, X.R. Li, G. Oxholm, T. Gittings, T. Bui, N. Flammarion, J. Collomosse. *ARIA: Adversarially Robust Image Attribution for Content Provenance* (CVPR 2022 Workshop on Media Forensics)

F. Croce, **M. Andriushchenko**, N. Singh, N. Flammarion, M. Hein. *Sparse-RS: a versatile framework for query-efficient sparse black-box adversarial attacks* (AAAI 2022)

F. Croce*, **M. Andriushchenko***, V. Sehwag*, E. Debenedetti*, N. Flammarion, M. Chiang, P. Mittal, M. Hein. *RobustBench: a standardized adversarial robustness benchmark* (NeurIPS 2021 Datasets and Benchmarks Track, **Best Paper Honorable Mention Prize** at ICLR 2021 Workshop on Security and Safety in Machine Learning Systems)

M. Mosbach, **M. Andriushchenko**, D. Klakow. *On the Stability of Fine-tuning BERT: Misconceptions, Explanations, and Strong Baselines* (ICLR 2021)

M. Andriushchenko*, F. Croce*, N. Flammarion, M. Hein. *Square Attack: a Query-Efficient Black-Box Adversarial Attack via Random Search* (ECCV 2020)

M. Andriushchenko, N. Flammarion. *Understanding and Improving Fast Adversarial Training* (NeurIPS 2020)

M. Andriushchenko, M. Hein. *Provably Robust Boosted Decision Stumps and Trees against Adversarial Attacks* (NeurIPS 2019, contributed talk at Workshop on Machine Learning with Guarantees; **best paper award** at Swiss Machine Learning Day (2019))

M. Hein, **M. Andriushchenko**, J. Bitterwolf. *Why ReLU networks yield high-confidence predictions far away from the training data and how to mitigate the problem* (**oral** at CVPR 2019, 5.6% acceptance rate, contributed talk at ICML 2019 Uncertainty and Robustness in Deep Learning Workshop)

F. Croce*, **M. Andriushchenko***, M. Hein. *Provable Robustness of ReLU Networks via Maximization of Linear Regions* (AISTATS 2019)

M. Mosbach*, **M. Andriushchenko***, T. Trost, M. Hein, D. Klakow. *Logit Pairing Methods Can Fool Gradient-Based Attacks* (NeurIPS 2018 Workshop on Security in ML)

M. Hein and **M. Andriushchenko**. *Formal Guarantees on the Robustness of a Classifier Against Adversarial Manipulation* (NeurIPS 2017)

References

1. **Prof. Nicolas Flammarion**
   École Polytechnique Fédérale de Lausanne
   EPFL Station 14, 1015 Lausanne,
   Switzerland
   nicolas.flammarion@epfl.ch

2. **Prof. Matthias Hein**
   University of Tübingen
   Maria-von-Linden-Str. 6, 72076
   Tübingen, Germany
   matthias.hein@uni-tuebingen.de

3. **Prof. Zico Kolter**
   Carnegie Mellon University
   5000 Forbes Avenue, Pittsburgh,
   PA 15213, USA
   zkolter@cs.cmu.edu

4. **Prof. Eric Wong**
   University of Pennsylvania
   3330 Walnut St, Philadelphia,
   PA 19104, USA
   exwong@cis.upenn.edu

5. **Prof. Aleksander Mądry**
   Massachusetts Institute of Technology
   32 Vassar Street, Cambridge,
   MA 02139, USA
   madry.letters@mit.edu