

Maksym Andriushchenko – Curriculum Vitae

General Information

Site: <https://andriushchenko.me/>
Email: maksym@andriushchenko.me

Scholar: <https://scholar.google.com/citations?user=ZNtuJYoAAAAJ>
Github: <https://github.com/max-andr/>

Education

École Polytechnique Fédérale de Lausanne (EPFL), Switzerland (Sep 2019 - Oct 2024)

PhD Degree in Computer Science advised by Nicolas Flammarion
Thesis title: “Understanding Generalization and Robustness in Modern Deep Learning”
Committee: Zico Kolter, Sébastien Bubeck, Rachid Guerraoui, Florent Krzakala

Saarland University, Germany (Oct 2016 – Aug 2019)

Master’s Degree in Computer Science advised by Matthias Hein from the University of Tübingen
Thesis title: “Provable Adversarial Defenses for Boosting”

Dnipropetrovsk National University of Railway Transport, Ukraine (Sep 2012 – June 2016)

Bachelor’s Degree in Software Engineering — *with honors*

Work

ELLIS Institute and Max Planck Institute	Time: September 2025 – now Role: Principal investigator at the ELLIS Institute Tübingen and independent group leader at the Max Planck Institute for Intelligent Systems
EPFL	Time: September 2019 – August 2025 Role: PhD student → postdoctoral researcher supervised by Nicolas Flammarion.
Gray Swan AI	Time: February 2024 – November 2024 (part-time consultant) Role: Member of technical staff working with leading AI safety organizations (Anthropic, UK AI Safety Institute, Center for AI Safety)
Adobe Research	Time: July 2021 – October 2021 Role: Research internship supervised by John Collomosse
PrivatBank	Time: November 2015 – June 2016 (part-time) Role: Data scientist working on predictive modeling and e-commerce applications

Awards and Grants

Main awards	ELLIS PhD Award (\$2.5k award) Patrick Denantes Memorial Prize for the best PhD thesis in the School of Computer and Communication Sciences at EPFL (\$5k award) G-Research PhD Thesis Prize in Quantitative Research (2nd place, \$3k award) Google PhD fellowship 2022-2025 (\$80k per year for 3 years) Open Philanthropy AI PhD Fellowship 2022-2024 (\$10k per year for travel/equipment) EDIC PhD fellowship from EPFL for the first year (\$60k) DAAD MSc scholarship for 2 years to study at Saarland University (\$20k)
Independent grants	Schmidt Sciences, Safety in the Inference-Time Compute Paradigm (co-PI; in processing) Academic Start-up Package, Technical AI Safety Program by Open Philanthropy (in processing)
Co-written grants	AI Safety Science grant (\$200k funded by Schmidt Sciences in 2025) Safe GenAI via Robust Content Moderation Models (\$100k funded by Google in 2024) Safety, Robustness, and Alignment of LLM agents (\$40k funded by Google in 2024) Google Research Collab 2022-2023 (\$80k for one year + \$20k in cloud compute)
Awards for papers and competitions	Trojan detection challenge at SatML'24: first place Swiss AI Safety Prize (2024): award for one of the top paper submissions Joint Conference of Korean AI Association (2023): best paper award ICLR Workshop on Security & Safety in ML Systems (2021): best paper honorable mention Swiss Machine Learning Day (2019): best paper award

Press Coverage

2025	Forbes : Embedding LLM Circuit Breakers Into AI Might Save Us From A Whole Lot Of Ghastly Troubles Le Figaro : When ChatGPT Goes Off the Rails, Should We Be Worried? 24 heures : When ChatGPT Goes Off the Rails, Should We Be Worried?
2024	SwissInfo : Swiss researchers find security flaws in AI models Le Temps : Bomb making, cannibalism or cyber attacks: it is always possible to ask AI for advice on explosive topics Blick : EPFL discovers security flaws in several AI models TechXplore : Can we convince AI to answer harmful requests? Léman Bleu : EPFL: Security flaws in AI models ActuIA : EPFL study: the limits of LLMs in the face of adaptive attacks

Academic Service

Participant	International AI Safety Report (Second Edition): Chapter Lead (July 2025 - February 2026) Wilson Center Fellowship Program: A Path Toward AI Policy (November 2024 - January 2025) Red teaming of the OpenAI fine-tuning service as an external expert (October 2023) Robust AI 4-day workshop organized by Airbus AI Research and TNO (January 2021)
Area chair	ICML'26
Reviewer	PNAS, TMLR, ICLR'25, NeurIPS'24, NeurIPS'23, ICML'23, NeurIPS'22 (top reviewer), ICML'22, NeurIPS'21, ICML'21, CVPR'21, ICLR'21 (outstanding reviewer), NeurIPS'20 (top 10% reviewers)
Program committee in workshops	NeurIPS'24 Red Teaming GenAI Workshop, NeurIPS'24 SATA Workshop: Towards Safe & Trustworthy Agents, NeurIPS'24 3rd Workshop on New Frontiers in Adversarial ML, ICML'24 Workshop on the Next Generation of AI Safety, NeurIPS'23 R0-FoMo Workshop on Robustness of Few-shot and Zero-shot Learning in Foundation Models, NeurIPS'23 Workshop on Distribution Shifts: New Frontiers with Foundation Models, ICML'23 2nd ICML Workshop on New Frontiers in Adversarial ML, ICLR'23 Workshop on Pitfalls of Limited Data and Computation for Trustworthy ML, NeurIPS'22 Workshop on Distribution Shifts, NeurIPS'22 ML Safety Workshop, ICML'22 New Frontiers in Adversarial ML, ICML'22 Principles of Distribution Shift, NeurIPS'21 Distribution Shifts: Connecting Methods and Applications, ICML'21 Uncertainty and Robustness in Deep Learning, CVPR'21 Adversarial ML in Real-World Computer Vision Systems, ICLR'21 Robust and Reliable ML in the Real World, Security and Safety in ML Systems, ICML'20 Uncertainty and Robustness in Deep Learning, CVPR'20 Adversarial ML in Computer Vision, ICLR'20 Towards Trustworthy ML (best reviewer award)
Outreach activities	National coordinator for Switzerland at #ScienceForUkraine Coordinator for Switzerland and admission officer at the Ukrainian Global University AI lectures at a summer camp for displaced Ukrainian children in Romania

Student Supervision

Agatha Duzan (EPFL)	MSc Project (2025) : A Benchmark for Measuring Safety of Computer Use Agents
Thomas Kuntz (EPFL)	MSc Project (2025) : A Benchmark for Measuring Safety of Computer Use Agents
Alexander Panfilov (University of Tübingen)	PhD thesis (2024-2028) : A Jailbreaking Perspective on LLM Safety (co-supervised with Jonas Geiping)
Joshua Freeman (ETH Zürich)	MSc project (2024) : Exploring Memorization and Copyright Violation in Frontier Large Language Models (<i>accepted at NeurIPS 2024 Safe Generative AI Workshop</i>)
Hao Zhao (EPFL)	MSc thesis (2023) : Long Is More for Alignment: A Simple but Tough-to-Beat Baseline for Instruction Fine-Tuning (<i>published at ICML'24, nominated for EPFL Outstanding Master's Thesis, featured in MIT Technology Review China</i>)
Hichem Hadhri (EPFL)	MSc project (2023) : Understanding overfitting in large language models
Tiberiu Musat (EPFL)	BSc project (2023) : Investigating key components for fast optimization of deep networks

Francesco d'Angelo (EPFL)	PhD semester project (2023): Understanding the role of weight decay in deep learning (<i>published at NeurIPS'24</i>)
Théau Vannier (EPFL)	MSc project (2023): Understanding the training instability of transformers
Joshua Freeman (EPFL)	BSc project (2022, unofficial): Automatic recognition of unexploded ordnance using transfer learning
Jana Vuckovic (EPFL)	MSc project (2022): Rethinking the relationship between sharpness and generalization (<i>follow-up work is published at ICML'23</i>)
Mehrdad Saberi (EPFL)	Summer internship (2021): Wasserstein adversarial training and perceptual robustness
Edoardo Debenedetti (EPFL)	MSc project (2021): RobustBench: a standardized adversarial robustness benchmark (<i>published at NeurIPS'21 Datasets and Benchmarks Track; Best Paper Honorable Mention Prize at ICLR 2021 Workshop on Security and Safety in ML Systems</i>)
Klim Kireev (EPFL)	PhD semester project (2020): On the effectiveness of adversarial training against common corruptions (<i>published at UAI'22</i>)
Etienne Bonvin (EPFL)	MSc project (2020): Adversarial robustness of kernel methods
Oriol Barbany (EPFL)	MSc project (2019): Affine-invariant robust training (co-advised with Sebastian Stich)

Teaching Experience

EPFL	Probability & Statistics 2021, 2022 (by E. Abbé), Machine Learning 2020, 2021, 2022, 2023 (by M. Jaggi, N. Flammarion), Advanced Algorithms 2020 (by M. Kapralov)
MPI for Informatics	Machine Learning 2018-2019 (lecturer: B. Schiele)
Saarland University	Neural Networks: Implementation and Application 2017 (lecturer: D. Klakow)

Selected Publications

- M. Andriushchenko**, A. Souly, M. Dziemian, D. Duenas, M. Lin, J. Wang, D. Hendrycks, A. Zou, Z. Kolter, M. Fredrikson, E. Winsor, J. Wynne, Y. Gal, X. Davies. AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents ([ICLR 2025](#))
- M. Andriushchenko**, F. Croce, N. Flammarion. Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks ([ICLR 2025](#))
- M. Andriushchenko**, N. Flammarion. Towards Understanding Sharpness-Aware Minimization ([ICML 2022](#))
- F. Croce*, **M. Andriushchenko***, V. Sehwag*, E. Debenedetti*, N. Flammarion, M. Chiang, P. Mittal, M. Hein. RobustBench: a standardized adversarial robustness benchmark ([NeurIPS 2021 Datasets and Benchmarks Track](#), **Best Paper Honorable Mention Prize at ICLR 2021 Workshop on Security and Safety in Machine Learning Systems**)
- M. Andriushchenko***, F. Croce*, N. Flammarion, M. Hein. Square Attack: a Query-Efficient Black-Box Adversarial Attack via Random Search ([ECCV 2020](#))

Full Publication List

- T. Kuntz, A. Duzan, H. Zhao, F. Croce, Z. Kolter, N. Flammarion, **M. Andriushchenko**. OS-Harm: A Benchmark for Measuring Safety of Computer Use Agents ([arXiv, June 2025](#); **oral at ICML'25 WCUA workshop**)
- C. Yueh-Han, N. Joshi, Y. Chen, **M. Andriushchenko**, R. Angell, H. He. Monitoring Decomposition Attacks in LLMs with Lightweight Sequential Monitors ([arXiv, June 2025](#))
- A. Panfilov, P. Kassianik, **M. Andriushchenko**, J. Geiping. Capability-Based Scaling Laws for LLM Red-Teaming ([arXiv, May 2025](#))
- M. Andriushchenko**, A. Souly, M. Dziemian, D. Duenas, M. Lin, J. Wang, D. Hendrycks, A. Zou, Z. Kolter, M. Fredrikson, E. Winsor, J. Wynne, Y. Gal, X. Davies. AgentHarm: A Benchmark for Measuring Harmfulness of LLM Agents ([ICLR 2025](#))
- M. Andriushchenko**, F. Croce, N. Flammarion. Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks ([ICLR 2025](#))
- M. Andriushchenko**, N. Flammarion. Does Refusal Training in LLMs Generalize to the Past Tense? ([ICLR 2025](#))

H. Zhao, **M. Andriushchenko**, F. Croce, N. Flammarion. Is In-Context Learning Sufficient for Instruction Following in LLMs? ([ICLR 2025](#))

J. Freeman, C. Rippe, E. Debenedetti, **M. Andriushchenko**. Exploring Memorization and Copyright Violation in Frontier LLMs: A Study of the New York Times v. OpenAI 2023 lawsuit ([NeurIPS 2024 Safe Generative AI Workshop](#))

A. Zou, L. Phan, J. Wang, D. Duenas, M. Lin, **M. Andriushchenko**, R. Wang, Z. Kolter, M. Fredrikson, D. Hendrycks. Improving Alignment and Robustness with Short Circuiting ([NeurIPS 2024](#))

M. Andriushchenko*, F. D'Angelo*, A. Varre, N. Flammarion. Why Do We Need Weight Decay in Modern Deep Learning? ([NeurIPS 2024](#))

P. Chao*, E. Debenedetti*, A. Robey*, **M. Andriushchenko***, F. Croce, V. Sehwag, E. Dobriban, N. Flammarion, G.J. Pappas, F. Tramer, H. Hassani, E. Wong. JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models ([NeurIPS 2024 Datasets and Benchmarks Track](#))

J. Rando, F. Croce, K. Mitka, S. Shabalin, **M. Andriushchenko**, N. Flammarion, F. Tramèr. Competition Report: Finding Universal Jailbreak Backdoors in Aligned LLMs ([arXiv, April 2024](#))

H. Zhao, **M. Andriushchenko**, F. Croce, N. Flammarion. Long Is More for Alignment: A Simple but Tough-to-Beat Baseline for Instruction Fine-Tuning ([ICML 2024](#))

L. Adilova, **M. Andriushchenko**, M. Kamp, A. Fischer, M. Jaggi. Layer-Wise Linear Mode Connectivity ([ICLR 2024](#))

M. Andriushchenko. Adversarial Attacks on GPT-4 via Simple Random Search ([December 2023](#))

E. Debenedetti, Z. Wan, **M. Andriushchenko**, V. Sehwag, K. Bhardwaj, B. Kailkhura. Scaling Compute Is Not All You Need for Adversarial Robustness ([ICLR 2024 Workshop on Reliable and Responsible Foundation Models](#))

S. Shin, D. Lee, **M. Andriushchenko**, N. Lee. The Effects of Overparameterization on Sharpness-Aware Minimization: An Empirical and Theoretical Analysis ([UAI 2025](#), **best paper award** at the Joint Conference of Korean Artificial Intelligence Association (2023))

M. Andriushchenko, D. Bahri, H. Mobahi, N. Flammarion. Sharpness-Aware Minimization Leads to Low-Rank Features ([NeurIPS 2023](#))

K. Kireev, **M. Andriushchenko**, C. Troncoso, N. Flammarion. Transferable Adversarial Robustness for Categorical Data via Universal Robust Embeddings ([NeurIPS 2023](#))

M. Andriushchenko, F. Croce, M. Müller, M. Hein, N. Flammarion. A modern look at the relationship between sharpness and generalization. ([ICML 2023](#))

M. Andriushchenko, A. Varre, L. Pillaud-Vivien, N. Flammarion. SGD with large step sizes learns sparse features ([ICML 2023](#))

K. Kireev*, **M. Andriushchenko***, N. Flammarion. On the effectiveness of adversarial training against common corruptions ([UAI 2022](#))

Michael Rose, Sanita Reinsone, **Maksym Andriushchenko**, Marcin Bartosiak, Anna Bobak et al. #ScienceForUkraine: an Initiative to Support the Ukrainian Academic Community. “3 Months Since Russia’s Invasion in Ukraine”, February 26 – May 31, 2022 ([SSRN, 2022](#))

M. Andriushchenko, N. Flammarion. Towards Understanding Sharpness-Aware Minimization ([ICML 2022](#))

M. Andriushchenko, X. Rebecca Li, Geoffrey Oxholm, Thomas Gittings, Tu Bui, Nicolas Flammarion, John Collomosse. ARIA: Adversarially Robust Image Attribution for Content Provenance ([CVPR 2022 Workshop on Media Forensics](#))

F. Croce, **M. Andriushchenko**, N. Singh, N. Flammarion, M. Hein. Sparse-RS: a versatile framework for query-efficient sparse black-box adversarial attacks ([AAAI 2022](#))

F. Croce*, **M. Andriushchenko***, V. Sehwag*, E. Debenedetti*, N. Flammarion, M. Chiang, P. Mittal, M. Hein. RobustBench: a standardized adversarial robustness benchmark ([NeurIPS 2021 Datasets and Benchmarks Track](#), **Best Paper Honorable Mention Prize** at [ICLR 2021 Workshop on Security and Safety in Machine Learning Systems](#))

M. Mosbach, **M. Andriushchenko**, D. Klakow. On the Stability of Fine-tuning BERT: Misconceptions, Explanations, and Strong Baselines ([ICLR 2021](#))

M. Andriushchenko*, F. Croce*, N. Flammarion, M. Hein. Square Attack: a query-efficient black-box adversarial attack via random search ([ECCV 2020](#))

M. Andriushchenko, N. Flammarion. Understanding and Improving Fast Adversarial Training ([NeurIPS 2020](#))

M. Andriushchenko, M. Hein. Provably Robust Boosted Decision Stumps and Trees against Adversarial Attacks ([NeurIPS 2019](#), contributed talk at [Workshop on Machine Learning with Guarantees](#); **best paper award** at Swiss Machine Learning Day (2019))

M. Hein, **M. Andriushchenko**, J. Bitterwolf. Why ReLU networks yield high-confidence predictions far away from the training data and how to mitigate the problem (**oral** at [CVPR 2019](#), 5.6% acceptance rate, contributed talk at [ICML 2019 Uncertainty and Robustness in Deep Learning Workshop](#))

F. Croce*, **M. Andriushchenko***, M. Hein. Provable Robustness of ReLU Networks via Maximization of Linear Regions ([AISTATS 2019](#))

M. Mosbach*, **M. Andriushchenko***, T. Trost, M. Hein, D. Klakow. Logit Pairing Methods Can Fool Gradient-Based Attacks ([NeurIPS 2018 Workshop on Security in ML](#))

M. Hein and **M. Andriushchenko**. Formal Guarantees on the Robustness of a Classifier Against Adversarial Manipulation ([NeurIPS 2017](#))

References

1. Prof. Nicolas Flammarion

École Polytechnique Fédérale de Lausanne
EPFL Station 14, 1015 Lausanne,
Switzerland
nicolas.flammarion@epfl.ch

2. Prof. Matthias Hein

University of Tübingen
Maria-von-Linden-Str. 6, 72076
Tübingen, Germany
matthias.hein@uni-tuebingen.de

3. Prof. Zico Kolter

Carnegie Mellon University
5000 Forbes Avenue, Pittsburgh, PA 15213, USA
zkolter@cs.cmu.edu

4. Prof. Eric Wong

University of Pennsylvania
3330 Walnut St, Philadelphia,
PA 19104, USA
exwong@cis.upenn.edu

5. Prof. Aleksander Mądry

Massachusetts Institute of Technology
32 Vassar Street, Cambridge,
MA 02139, USA
madry.letters@mit.edu